

I used google benchmark for my measurements. I randomize the numbers and seed every time and after running each part several times I believe my report reflects the results on average.

```
user@DESKTOP-LAVV9R:~/assignments$ g++ -std=c++17 -O2 -I/usr/local/include -L/usr/local/lib task1.cpp -lbenchmark -lpthread -o sum
user@DESKTOP-LAVV9R:~/assignments$ ./sum
2025-11-23T21:09:06+01:00
Running ./sum
Run on (28 X 3417.6 MHz CPU s)
CPU Caches:
  L1 Data 48 KiB (x14)
  L1 Instruction 32 KiB (x14)
  L2 Unified 2048 KiB (x14)
  L3 Unified 33792 KiB (x1)
Load Average: 0.00, 0.00, 0.00
-----
Benchmark      Time          CPU   Iterations
Part1        220497 ns     230940 ns       3013
Part2        711523 ns     745223 ns       944
Part3        320395 ns     335573 ns       2088
Part4        315015 ns     304175 ns       2312
user@DESKTOP-LAVV9R:~/assignments$
```

Here is an image of all of the parts running and their results.

Part 1:

```
user@DESKTOP-LAVV9R:~/assignments$ g++ -std=c++17 -O2 -I/usr/local/include -L/usr/local/lib task1.cpp -lbenchmark -lpthread -o sum
user@DESKTOP-LAVV9R:~/assignments$ ./sum
2025-11-23T21:15:13+01:00
Running ./sum
Run on (28 X 3417.6 MHz CPU s)
CPU Caches:
  L1 Data 48 KiB (x14)
  L1 Instruction 32 KiB (x14)
  L2 Unified 2048 KiB (x14)
  L3 Unified 33792 KiB (x1)
Load Average: 0.04, 0.01, 0.00
-----
Benchmark      Time          CPU   Iterations
Part1        227368 ns     227374 ns       3119
user@DESKTOP-LAVV9R:~/assignments$
```

Part 2:

```
user@DESKTOP-LAVV9R:~/assignments$ g++ -std=c++17 -O2 -I/usr/local/include -L/usr/local/lib task1.cpp -lbenchmark -lpthread -o sum
user@DESKTOP-LAVV9R:~/assignments$ ./sum
2025-11-23T20:09:07+01:00
Running ./sum
Run on (28 X 3417.6 MHz CPU s)
CPU Caches:
  L1 Data 48 KiB (x14)
  L1 Instruction 32 KiB (x14)
  L2 Unified 2048 KiB (x14)
  L3 Unified 33792 KiB (x1)
Load Average: 0.14, 0.09, 0.02
-----
Benchmark      Time          CPU   Iterations
Part2        780285 ns     740894 ns       931
user@DESKTOP-LAVV9R:~/assignments$
```

The time and CPU usage for part 2 is around 3 times higher for both.

This makes sense as significantly more operations are performed which requires more CPU usage. As iterations are significantly lower than any other task, pipelining is possibly a large contributing factor to performance as it will allow for more instructions and calculations to be done per cycle. As I ran these tests on a powerful PC, the amount of cores and fast cache reduces the time it takes to complete instructions and will reduce the time the CPU waits for data to move between memory.

### Part 3:

```
user@DESKTOP-LAVVV9R:~/assignments$ g++ -std=c++17 -O2 -I/usr/local/include -L/usr/local/lib task1.cpp -lbenchmark -lpthread -o sum
user@DESKTOP-LAVVV9R:~/assignments$ ./sum
2025-11-23T20:22:02+01:00
Running ./sum
Run on (28 X 3417.6 MHz CPU s)
CPU Caches:
  L1 Data 48 KiB (x14)
  L1 Instruction 32 KiB (x14)
  L2 Unified 2048 KiB (x14)
  L3 Unified 33792 KiB (x1)
Load Average: 0.08, 0.05, 0.01
-----
Benchmark      Time       CPU Iterations
-----
Part3      311451 ns     331879 ns      2099
user@DESKTOP-LAVVV9R:~/assignments$
```

The performance of part 3 is a lot better than part 2. One reason it may be faster is because the values have to be accessed from the array 2 times as much in part 2 than part 3 which increases the time the CPU is waiting for values to transfer across memory and registers. The compiler also may be more optimized when everything is in one expression.

### Part 4:

```
user@DESKTOP-LAVVV9R:~/assignments$ g++ -std=c++17 -O2 -I/usr/local/include -L/usr/local/lib task1.cpp -lbenchmark -lpthread -o sum
user@DESKTOP-LAVVV9R:~/assignments$ ./sum
2025-11-23T20:36:23+01:00
Running ./sum
Run on (28 X 3417.6 MHz CPU s)
CPU Caches:
  L1 Data 48 KiB (x14)
  L1 Instruction 32 KiB (x14)
  L2 Unified 2048 KiB (x14)
  L3 Unified 33792 KiB (x1)
Load Average: 0.19, 0.17, 0.06
-----
Benchmark      Time       CPU Iterations
-----
Part4      245890 ns     257220 ns      2758
user@DESKTOP-LAVVV9R:~/assignments$
```

It doesn't affect throughput much. The interactions are lowered but CPU and Time actually increase. The if statement will make branches and any mis prediction of the branches will impact the performance. The slowdown may also be because not enough calculations are skipped as 500000 may be too high to impact the performance.